

Gradient Bootstrapping

Anushree Shivarudrappa

June 15, 2016

1. Pre-Processing

```
library(data.table)
library(ggplot2)
library(dplyr)
library(scales)
library(RColorBrewer)
library(tidyr)
library(caTools)
library(rpart)
library(rpart.plot)
library(ROCR)
library(randomForest)
library(tree)
library(caret)
library(e1071)
library(gbm)
```

2. Data Loading

```
Death_US <- fread("DeathRecords.csv", header = T)
```

3. Selecting dataset for model

```
# separates natural death
Death_US_natural <- Death_US[Death_US$MannerOfDeath == 7, ]
```

Select required variables

```
require(MASS)
require(dplyr)
natural_sub <- Death_US_natural %>% dplyr::select(Education2003Revision, Sex, Age,
  InfantAgeRecode22,
  PlaceOfDeathAndDecedentsStatus, MaritalStatus, InjuryAtWork,
  MannerOfDeath,
  Autopsy, ActivityCode, PlaceOfInjury, Icd10Code, CauseRecode358,
  CauseRecode113, InfantCauseRecode130, CauseRecode39,
  NumberOfEntityAxisConditions, NumberOfRecordAxisConditions, Race)
```

Converting Character variable into Integer variable

```
natural_sub$Sex <- as.integer(as.factor(natural_sub$Sex))
natural_sub$MaritalStatus <- as.integer(as.factor(natural_sub$MaritalStatus))
natural_sub$InjuryAtWork <- as.integer(as.factor(natural_sub$InjuryAtWork))
natural_sub$Autopsy <- gsub("n", "N", natural_sub$Autopsy)
natural_sub$Autopsy <- as.integer(as.factor(natural_sub$Autopsy))
natural_sub$Icd10Code <- as.integer(as.factor(natural_sub$Icd10Code))
```

As we analyzed the feature variables are “Age + InfantAgeRecode22 + PlaceOfDeathAndDecedentsStatus + MaritalStatus + ActivityCode + PlaceOfInjury + NumberOfRecordAxisConditions + NumberOfEntityAxisConditions”

```
# Since the decision tree support till 32 levels removing 7 levels which has less entries
table(factor(natural_sub$CauseRecode39))
```

```
##
##      1      2      3      5      6      7      8      9     10     11
##    366    37   5619   9053  43839  33847 133412  34621  23359  23422
##     12    13    14    15    16    17    20    21    22    23
##  25734 17116 19671 133276  63721  75552  37415 310848 175752  23704
##     24    25    26    27    28    29    30    31    32    33
## 111664   5426 16551  45801 125752   2519  31595  41369   1000   9930
##     34    35    36    37    38    39    40    41    42
##    8110   414  23035 433081   212  13088     8     5     9
```

```
CauseExtraRemove <- natural_sub[, natural_sub$CauseRecode39 %in% c(2, 40, 41, 42, 38, 35, 1)]
table(CauseExtraRemove)
```

```
## CauseExtraRemove
##   FALSE   TRUE
## 2058882   1051
```

```
# remove the 7 factors levels from Death_US_natural dataset
natural_sub <- natural_sub[!(CauseExtraRemove)]
nrow(natural_sub)
```

```
## [1] 2058882
```

```
# model data
modeldata <- natural_sub

# We will do a random 70:30 split in our data set (70% will be for training models,
# 30% to evaluate them)
set.seed(111)
# randomly pick 70% of the number of observations
index <- sample.split(modeldata$CauseRecode39, SplitRatio = 0.7)
# subset data to include only the elements in the index
train <- subset(modeldata, index==T)
nrow(train)
```

```
## [1] 1441215
```

```
# subset data to include all but the elements in the index
test <- subset(modeldata, index==F)
nrow(test)
```

```
## [1] 617667
```

```
# take a copy of ICD10Code of test set and remove the variable from test set
Cause39 <- test$CauseRecode39
test$CauseRecode39 <- NULL
```

Model 3:: Gradient Bootstrapping

```
gbm2 <- gbm(as.factor(CauseRecode39) ~ Age + InfantAgeRecode22 +
            PlaceOfDeathAndDecedentsStatus + MaritalStatus + ActivityCode +
            PlaceOfInjury + NumberOfRecordAxisConditions +
            NumberOfEntityAxisConditions,
            data = train,
            var.monotone=c(0,0,0,0,0,0,0,0),
            # +1: monotone increase,
            # 0: no monotone restrictions
            distribution="gaussian",      # bernoulli, adaboost, gaussian,
            # poisson, coxph, and quantile available
            n.trees=3000,                 # number of trees
            shrinkage=0.005,              # shrinkage or learning rate,
            # 0.001 to 0.1 usually work
            interaction.depth=3,           # 1: additive model, 2: two-way interactions, etc.
            bag.fraction = 0.5,           # subsampling fraction, 0.5 is probably best
            n.minobsinnode = 10,          # minimum total weight needed in each node
            cv.folds = 5,                 # do 5-fold cross-validation
            keep.data=TRUE,               # keep a copy of the dataset with the object
            verbose=T )
```

## Iter	TrainDeviance	ValidDeviance	StepSize	Improve
## 1	75.6015	nan	0.0050	0.0400
## 2	75.5615	nan	0.0050	0.0397
## 3	75.5224	nan	0.0050	0.0393
## 4	75.4835	nan	0.0050	0.0388
## 5	75.4448	nan	0.0050	0.0385
## 6	75.4066	nan	0.0050	0.0381
## 7	75.3689	nan	0.0050	0.0377
## 8	75.3313	nan	0.0050	0.0374
## 9	75.2943	nan	0.0050	0.0370
## 10	75.2577	nan	0.0050	0.0366
## 20	74.9083	nan	0.0050	0.0337
## 40	74.2900	nan	0.0050	0.0288
## 60	73.7363	nan	0.0050	0.0264
## 80	73.2566	nan	0.0050	0.0234
## 100	72.8224	nan	0.0050	0.0192


##	120	72.4316	nan	0.0050	0.0174
##	140	72.0683	nan	0.0050	0.0157
##	160	71.7419	nan	0.0050	0.0141
##	180	71.4435	nan	0.0050	0.0139
##	200	71.1583	nan	0.0050	0.0147
##	220	70.9027	nan	0.0050	0.0107
##	240	70.6670	nan	0.0050	0.0124
##	260	70.4454	nan	0.0050	0.0096
##	280	70.2388	nan	0.0050	0.0083
##	300	70.0459	nan	0.0050	0.0077
##	320	69.8689	nan	0.0050	0.0096
##	340	69.7060	nan	0.0050	0.0070
##	360	69.5513	nan	0.0050	0.0075
##	380	69.4165	nan	0.0050	0.0055
##	400	69.2881	nan	0.0050	0.0051
##	420	69.1719	nan	0.0050	0.0058
##	440	69.0600	nan	0.0050	0.0062
##	460	68.9597	nan	0.0050	0.0039
##	480	68.8681	nan	0.0050	0.0053
##	500	68.7815	nan	0.0050	0.0034
##	520	68.7014	nan	0.0050	0.0032
##	540	68.6258	nan	0.0050	0.0030
##	560	68.5593	nan	0.0050	0.0039
##	580	68.4960	nan	0.0050	0.0031
##	600	68.4365	nan	0.0050	0.0034
##	620	68.3792	nan	0.0050	0.0032
##	640	68.3287	nan	0.0050	0.0021
##	660	68.2821	nan	0.0050	0.0019
##	680	68.2350	nan	0.0050	0.0019
##	700	68.1918	nan	0.0050	0.0021
##	720	68.1525	nan	0.0050	0.0017
##	740	68.1167	nan	0.0050	0.0016
##	760	68.0812	nan	0.0050	0.0015
##	780	68.0480	nan	0.0050	0.0016
##	800	68.0167	nan	0.0050	0.0016
##	820	67.9875	nan	0.0050	0.0011
##	840	67.9613	nan	0.0050	0.0012
##	860	67.9357	nan	0.0050	0.0012
##	880	67.9116	nan	0.0050	0.0013
##	900	67.8900	nan	0.0050	0.0011
##	920	67.8688	nan	0.0050	0.0011
##	940	67.8505	nan	0.0050	0.0008
##	960	67.8318	nan	0.0050	0.0008
##	980	67.8143	nan	0.0050	0.0007
##	1000	67.7978	nan	0.0050	0.0007
##	1020	67.7821	nan	0.0050	0.0009
##	1040	67.7674	nan	0.0050	0.0006
##	1060	67.7533	nan	0.0050	0.0008
##	1080	67.7399	nan	0.0050	0.0005
##	1100	67.7267	nan	0.0050	0.0007
##	1120	67.7132	nan	0.0050	0.0007
##	1140	67.7016	nan	0.0050	0.0004
##	1160	67.6909	nan	0.0050	0.0004
##	1180	67.6793	nan	0.0050	0.0005

##	1200	67.6696	nan	0.0050	0.0006
##	1220	67.6598	nan	0.0050	0.0004
##	1240	67.6504	nan	0.0050	0.0003
##	1260	67.6417	nan	0.0050	0.0003
##	1280	67.6327	nan	0.0050	0.0004
##	1300	67.6241	nan	0.0050	0.0003
##	1320	67.6158	nan	0.0050	0.0003
##	1340	67.6081	nan	0.0050	0.0003
##	1360	67.6004	nan	0.0050	0.0005
##	1380	67.5933	nan	0.0050	0.0005
##	1400	67.5857	nan	0.0050	0.0003
##	1420	67.5786	nan	0.0050	0.0002
##	1440	67.5720	nan	0.0050	0.0003
##	1460	67.5659	nan	0.0050	0.0004
##	1480	67.5598	nan	0.0050	0.0002
##	1500	67.5533	nan	0.0050	0.0003
##	1520	67.5472	nan	0.0050	0.0003
##	1540	67.5418	nan	0.0050	0.0002
##	1560	67.5361	nan	0.0050	0.0002
##	1580	67.5305	nan	0.0050	0.0003
##	1600	67.5252	nan	0.0050	0.0002
##	1620	67.5199	nan	0.0050	0.0003
##	1640	67.5152	nan	0.0050	0.0002
##	1660	67.5106	nan	0.0050	0.0002
##	1680	67.5056	nan	0.0050	0.0003
##	1700	67.5008	nan	0.0050	0.0002
##	1720	67.4962	nan	0.0050	0.0002
##	1740	67.4921	nan	0.0050	0.0002
##	1760	67.4875	nan	0.0050	0.0003
##	1780	67.4831	nan	0.0050	0.0002
##	1800	67.4785	nan	0.0050	0.0003
##	1820	67.4748	nan	0.0050	0.0001
##	1840	67.4708	nan	0.0050	0.0002
##	1860	67.4666	nan	0.0050	0.0003
##	1880	67.4632	nan	0.0050	0.0001
##	1900	67.4588	nan	0.0050	0.0002
##	1920	67.4551	nan	0.0050	0.0002
##	1940	67.4518	nan	0.0050	0.0001
##	1960	67.4479	nan	0.0050	0.0002
##	1980	67.4444	nan	0.0050	0.0002
##	2000	67.4406	nan	0.0050	0.0002
##	2020	67.4368	nan	0.0050	0.0001
##	2040	67.4333	nan	0.0050	0.0001
##	2060	67.4301	nan	0.0050	0.0001
##	2080	67.4273	nan	0.0050	0.0001
##	2100	67.4234	nan	0.0050	0.0002
##	2120	67.4200	nan	0.0050	0.0001
##	2140	67.4164	nan	0.0050	0.0001
##	2160	67.4136	nan	0.0050	0.0001
##	2180	67.4101	nan	0.0050	0.0001
##	2200	67.4068	nan	0.0050	0.0001
##	2220	67.4035	nan	0.0050	0.0001
##	2240	67.4003	nan	0.0050	0.0002
##	2260	67.3964	nan	0.0050	0.0001

##	2280	67.3934	nan	0.0050	0.0003
##	2300	67.3902	nan	0.0050	0.0001
##	2320	67.3872	nan	0.0050	0.0003
##	2340	67.3842	nan	0.0050	0.0002
##	2360	67.3813	nan	0.0050	0.0001
##	2380	67.3778	nan	0.0050	0.0002
##	2400	67.3741	nan	0.0050	0.0001
##	2420	67.3713	nan	0.0050	0.0002
##	2440	67.3685	nan	0.0050	0.0001
##	2460	67.3660	nan	0.0050	0.0001
##	2480	67.3631	nan	0.0050	0.0001
##	2500	67.3604	nan	0.0050	0.0001
##	2520	67.3573	nan	0.0050	0.0000
##	2540	67.3545	nan	0.0050	0.0001
##	2560	67.3519	nan	0.0050	0.0001
##	2580	67.3487	nan	0.0050	0.0002
##	2600	67.3456	nan	0.0050	0.0002
##	2620	67.3432	nan	0.0050	0.0001
##	2640	67.3407	nan	0.0050	0.0001
##	2660	67.3389	nan	0.0050	0.0001
##	2680	67.3363	nan	0.0050	0.0002
##	2700	67.3341	nan	0.0050	0.0001
##	2720	67.3319	nan	0.0050	0.0001
##	2740	67.3296	nan	0.0050	0.0001
##	2760	67.3273	nan	0.0050	0.0001
##	2780	67.3252	nan	0.0050	0.0000
##	2800	67.3231	nan	0.0050	0.0001
##	2820	67.3206	nan	0.0050	0.0001
##	2840	67.3188	nan	0.0050	0.0000
##	2860	67.3168	nan	0.0050	0.0000
##	2880	67.3148	nan	0.0050	0.0001
##	2900	67.3125	nan	0.0050	0.0000
##	2920	67.3107	nan	0.0050	0.0001
##	2940	67.3094	nan	0.0050	0.0000
##	2960	67.3077	nan	0.0050	0.0001
##	2980	67.3057	nan	0.0050	0.0000
##	3000	67.3040	nan	0.0050	0.0001

```
# check performance using an out-of-bag estimator
# OOB underestimates the optimal number of iterations
best.iter <- gbm.perf(gbm2,method="OOB")
```

```
## Warning in gbm.perf(gbm2, method = "OOB"): OOB generally underestimates the
## optimal number of iterations although predictive performance is reasonably
## competitive. Using cv.folds>0 when calling gbm usually results in improved
## predictive performance.
```



```
print(best.iter)
```

```
## [1] 3000
```

```
data.predict = predict(gbm2, n.trees = best.iter, newdata = test)
# Confusion matrix
conf_matrix <- table(data.predict, Cause39)
```

Accuracy of model and SSE

```
#Accuracy
sum(diag(conf_matrix)) / nrow(test)
```

```
## [1] 8.094977e-05
```

```
# SSE
SSE = sum((Cause39 - data.predict)^2)
print(SSE)
```

```
## [1] 65365605
```